

Stacked LLM defenses share a blind spot


what layering assumes

each layer is blind somewhere different



what we measured

the layers are blind in the same place



Every layer reads the same model, so no amount of diversity between layers removes the failure they have in common.