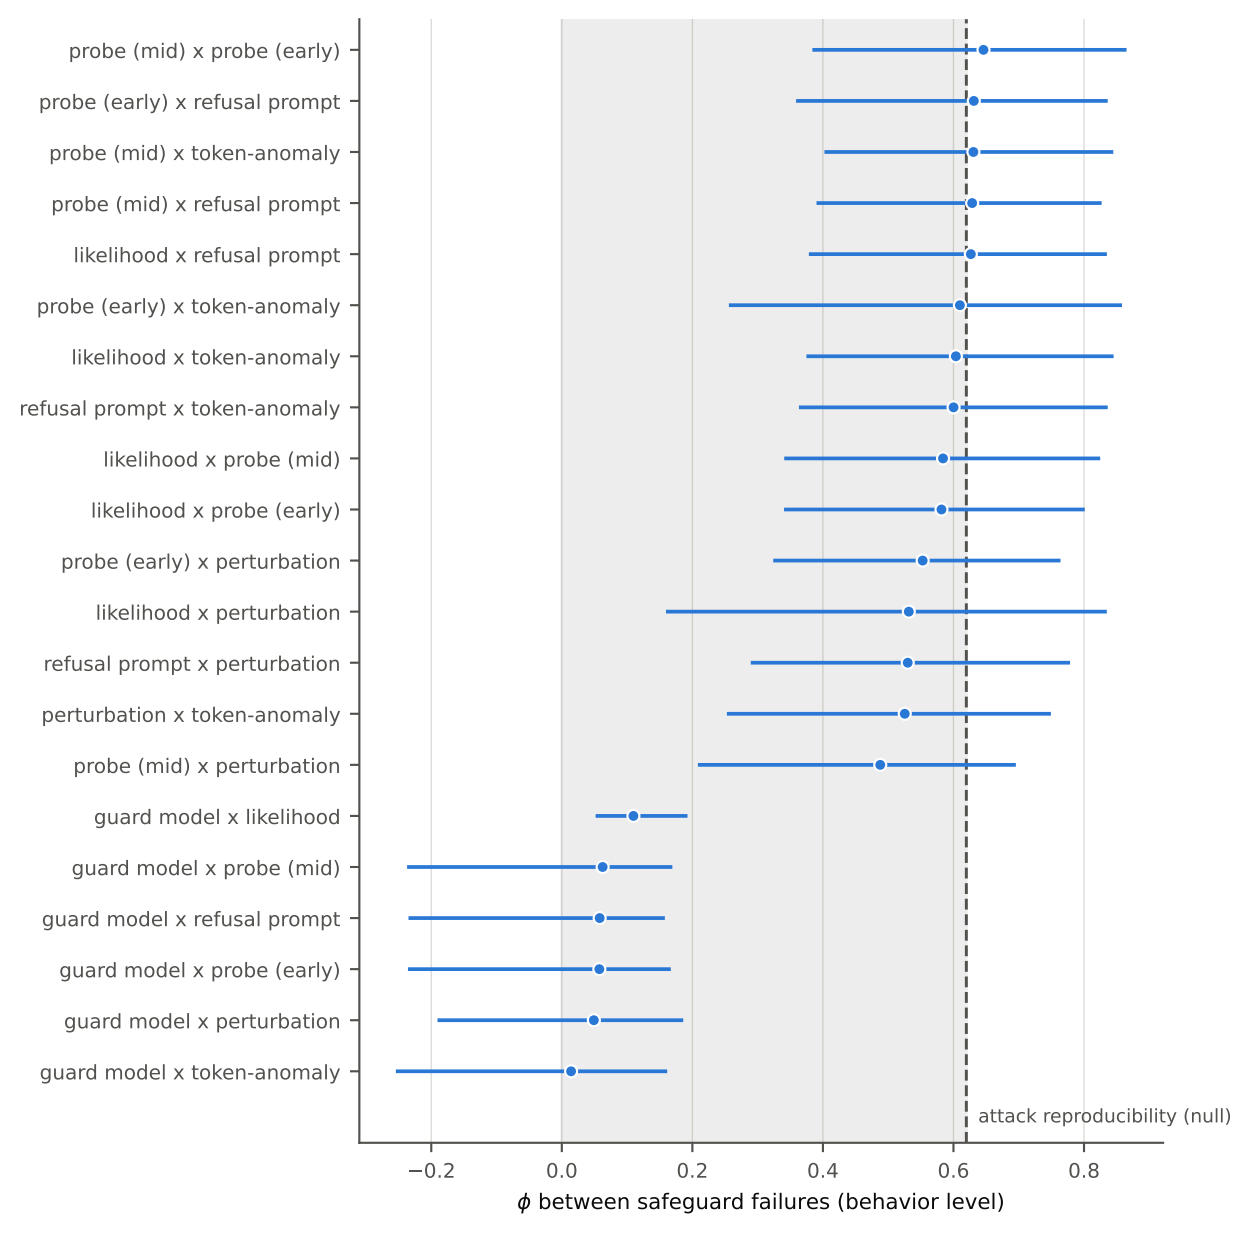
probe (mid) x probe (early)
probe (early) x refusal prompt
probe (mid) x token-anomaly
probe (mid) x refusal prompt
likelihood x refusal prompt
probe (early) x token-anomaly
likelihood x token-anomaly
refusal prompt x token-anomaly
likelihood x probe (mid)
likelihood x probe (early)
probe (early) x perturbation
likelihood x perturbation
refusal prompt x perturbation
perturbation x token-anomaly
probe (mid) x perturbation
guard model x likelihood
guard model x probe (mid)
guard model x refusal prompt
guard model x probe (early)
guard model x perturbation
guard model x token-anomaly
attack reproducibility (null)
−0.2
0.0
0.2
0.4
0.6
0.8
ϕ between safeguard failures (behavior level)