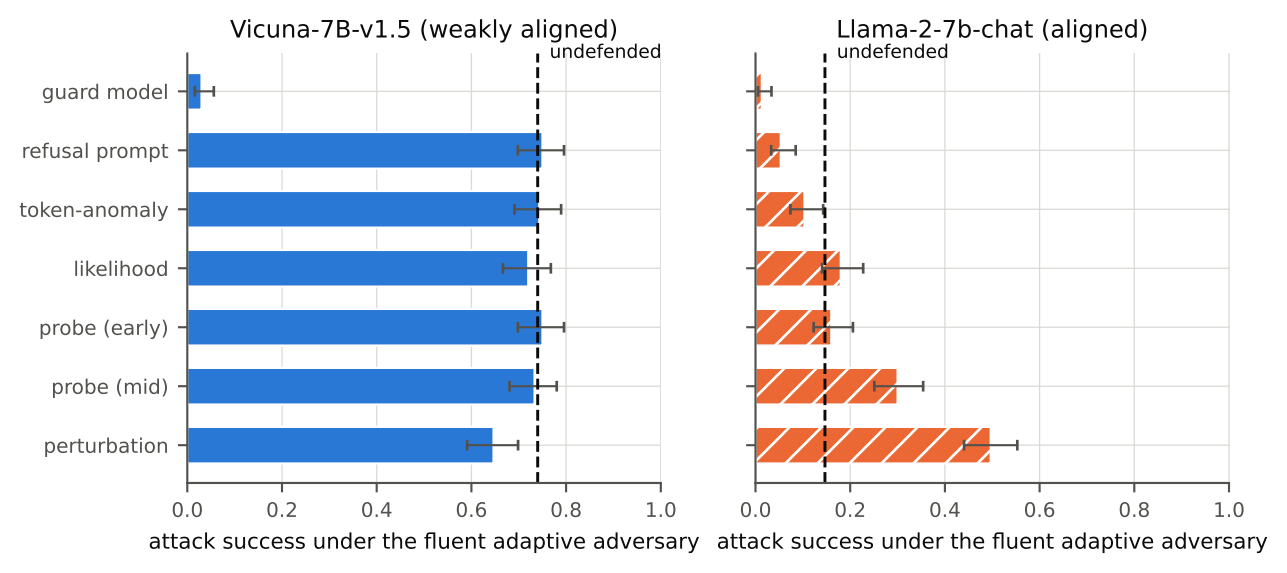
Vicuna-7B-v1.5 (weakly aligned)
undefended
guard model
refusal prompt
token-anomaly
likelihood
probe (early)
probe (mid)
perturbation
0.0
0.2
0.4
0.6
0.8
1.0
attack success under the fluent adaptive adversary
Llama-2-7b-chat (aligned)
undefended
0.0
0.2
0.4
0.6
0.8
1.0
attack success under the fluent adaptive adversary